

Project 1: Basic classification

WU1

By computing $\text{mean}((\text{datasets.TennisData.Y} > 0) == (h.\text{predictAll}(\text{datasets.TennisData.X}) > 0))$, you are actually computing $\frac{1}{N} \sum_{k=0}^{N-1} [y_k = f(\mathbf{x}_k)]$, which, by definition, is accuracy (1 - training error).

WU2

The training accuracy goes down after increasing number of examples. This is because there would be more ambiguity, meaning there would be examples with different labels with same feature values at non-leaf nodes where the DT makes a decision.

The testing error could go down even after you increase the number of examples. This is because the features that are considered important on training phase can vary as the number of examples changes. Let's say, when there are 50 examples, a feature x_{a1} is considered more dominant than the other features. However, as the number of examples increases, say 100, another feature x_{a2} may be considered more influential than x_{a1} . This may be right or wrong on the test data. If it is right on test data, the test accuracy goes up, and vice versa.

WU3

Please refer to the figure 1.

The training accuracy is guaranteed to monotonically increase, because it is table-look-up, except for the case which at every node on DT the probability of going right or left is 50-50. This will never increase the accuracy, but it is very unlikely to happen.

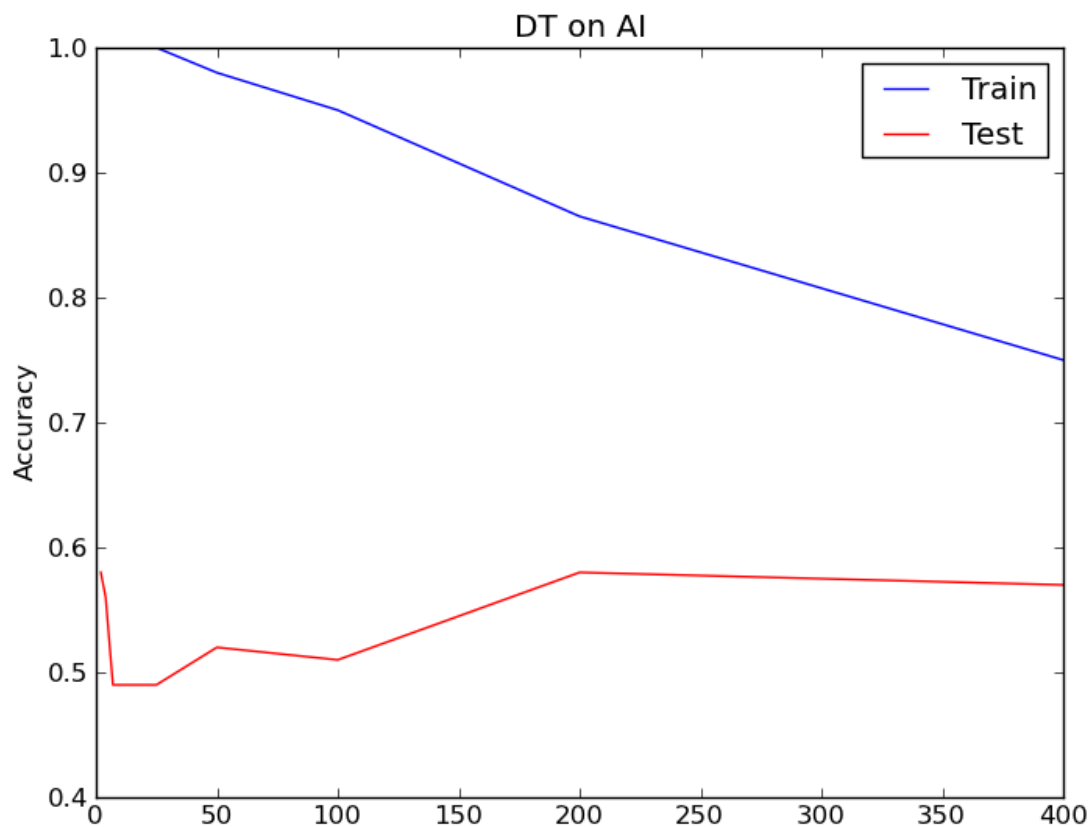
The testing accuracy could oscillate like on the figure, but not guaranteed. As questions asked at DT nodes go up, the model tends to overfit. This could either increase or decrease in accuracy. However, in the end the accuracy goes down overall.

WU4

Please refer to the figure 2.

```
-introduction to low-level programming concepts [212]
- program analysis and understanding [631]
- computer processing of pictorial information [733]
  - Leaf -1.0
  - Leaf 1.0
- introduction to human-computer interaction [434]
  - Leaf -1.0
  - Leaf 1.0
- computational linguistics ii [773]
```

Figure 1: Plot of DT with various numbers of examples



- computational methods [460]
 - Leaf -1.0
 - Leaf 1.0
- computational geometry [754]
 - Leaf 1.0
 - Leaf 1.0

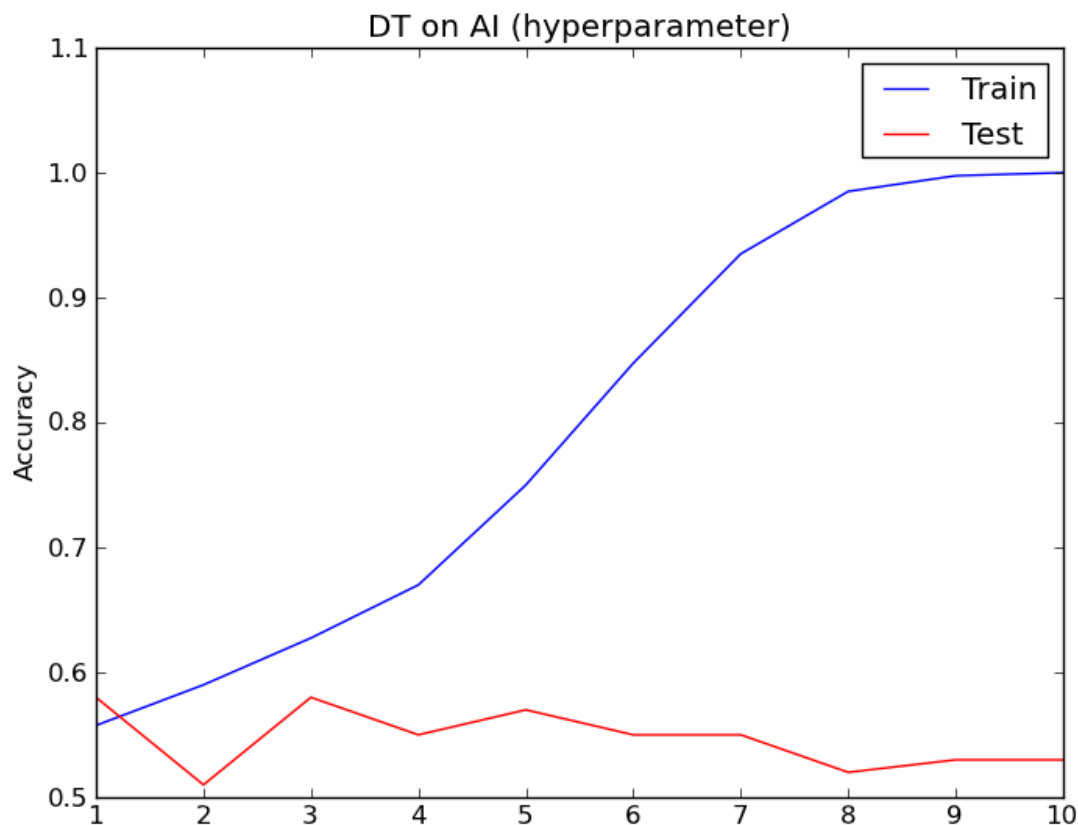
I may be wrong and there may be correlation between taking CG and taking these courses, but the courses look fairly random to me. This implies DT is not the best model for this problem.

WU5

For the course recommender data, generate train/test curves for varying values of K and epsilon (you figure out what are good ranges, this time). Include those curves: do you see evidence of overfitting and underfitting? Next, using $K=5$, generate learning curves for this data.

Figure 3 shows the training and test plots of KNN with varying k -values. As k approaches N , we are prone to underfitting, since our predictions are more and more based off the average of the examples in the training data. With $k = 1$, we are overfitting possible noise: for example, in the case where a data point we are asked

Figure 2: Plot of DT with various depth parameters



to predict lies in an area of mostly negatives, except for one positive that happens to be the nearest-neighbor.

Figure 4 shows the training and test plots of KNN with various ϵ -values. KNN with ϵ -ball nearest neighbors doesn't fit this problem well, only performing better-than-chance with $\epsilon = 5$. This could suggest that a lot of the data isn't necessarily close in the N -dimensional space.

Figure 5 shows the learning curve for KNN with $k = 5$. In general, we see an improvement in test accuracy as the number of training examples increases. This should be expected, since more training examples will increase the chances of KNN to be able to accurately predict at test time.

WU6

Your answer here.

Figure 3: Plot of KNN with various K-values

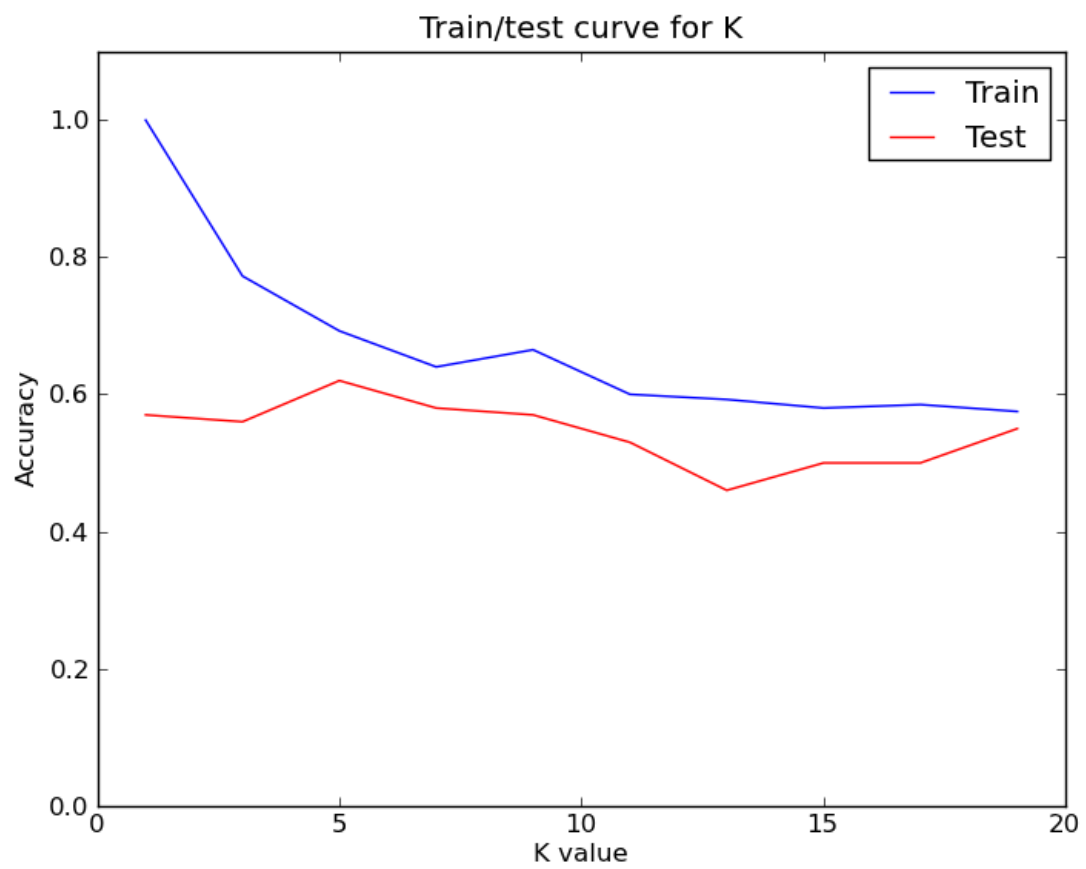


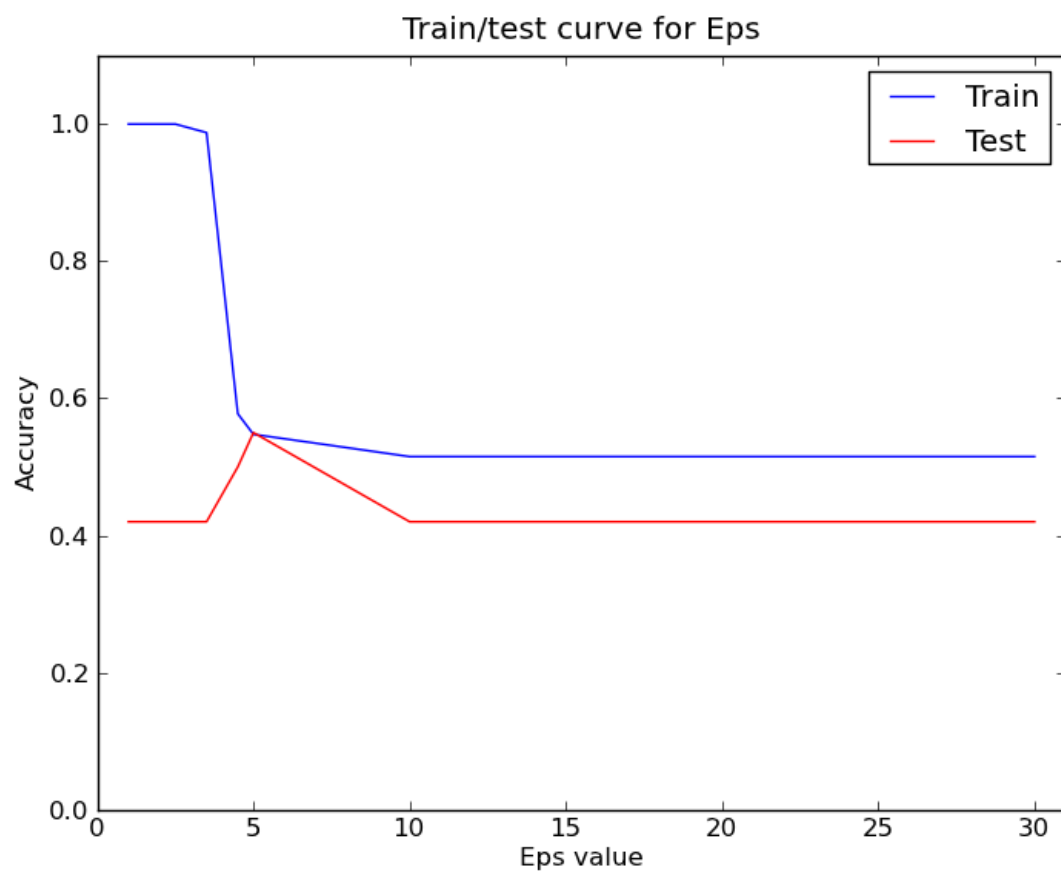
Figure 4: Plot of KNN with various ϵ -values

Figure 5: Learning Curve with K=5

